Team E: Responsible AI and governance (Teresa, Justyna)

Task 1 (PARALLEL after Definition Pack): Write Safety Boundaries (1 page)
Clear "never does" list:

- *not diagnosis*
- *not prescribing*
- *not replacing clinician*
- *always shows uncertainty and missing inputs*
- *always shows sources*

Task 2 (PARALLEL): Write Update and Non-Regression Policy
Simple gate:

- *new evidence can be added*
- *it must not break the vignette scores*
- *if it breaks them, it is flagged for review, not shipped as default*

This is your "don't get worse over time" story, in plain language.

TEAM E CONTRIBUTION BELOW ⬇️

In our MVP, safety is enforced by constraining system behavior rather than having the system make or suggest medical decisions. Safety boundaries are implemented through fixed vocabularies, structured inputs, explicit uncertainty language, and vignette-based evaluation, with human reviewers validating compliance.

1. Safety Boundaries (How the System Is Constrained)

We define explicit safety boundaries that limit what the system can and cannot do.
Each boundary is **implemented through concrete artifacts**, not abstract principles.

Key Safety Boundaries

1. **No Clinical Authority**

- The system does not recommend, rank, or prescribe treatments.
- Outputs describe evidence only, using neutral, descriptive language.

2. **No Guideline Interpretation**

- The system does not determine whether a treatment is “approved” or “allowed.”
- Evidence is labeled only as aligned, silent, or conflicting relative to guidelines.

3. **No Exhaustiveness Claims**

- The system never claims to show all relevant evidence.
- Uncertainty and incompleteness are explicitly stated.

4. **No Inference Beyond Inputs**

- The system uses only patient attributes provided in the vignette.
- Missing data leads to uncertainty, not assumptions.

5. **Strict Scope Control**

- Scope is fixed to NSCLC and treatment evidence only.
- No diagnostic, testing, or follow-up recommendations are produced.

How Boundaries Are Enforced

Each boundary is enforced via:

- Fixed output templates
- Restricted label vocabularies
- Structured input fields
- Yes/No evaluation checks in the scoring sheet

Safety boundary	**How it´s enforced**	**Where it´s visible**
Clinical authority	Descriptive-only output language	Evidence summaries, relevance scoring
Guideline interpretation	Restricted label set (aligned/silent/conflict)	Mapping_Accuracy sheet
Knowledge completeness	Mandatory uncertainty language	Output footer, one-pager
Input assumptions	No inference; explicit checks	vignette metadata sheet
Scope limitation	Treatment evidence only	Scoring sheet structure
Evaluation limits	Behaviour-based metrics only	Misclassification definition
Evaluation control	Vignette-based regression	Reused scoring sheet

Boundary compliance is verified through human review using the scoring sheets; enforcement is therefore socio-technical, not fully automated.

---

2. Update & Non-Regression Policy (How the System Can Change)

As evidence and system logic evolve, we require that updates do not silently degrade previously validated behavior.

What Can Change

- Evidence corpus (e.g., new PubMed entries)
- Retrieval and mapping logic
- Output formatting and clarity

What Is Frozen for MVP

- Scope (NSCLC, treatment evidence only)
- Input fields and vignette structure
- Label vocabulary
- Safety language and disclaimers

Or put in a table:

Allowed to Change (with controls)

Component	Examples
Evidence corpus	New PubMed entries
Parsing & retrieval logic	Improved keyword extraction
Mapping heuristics	Better guideline alignment rules
Output formatting	Clearer summaries, better UX

Frozen for MVP (no change without re-approval)

Component	Reason
Scope (NSCLC, treatment evidence only)	Prevent scope creep
Input fields	Maintain vignette comparability

Label vocabulary	Ensure consistent evaluation
Safety language templates	Prevent clinical overreach

---

3. Non-Regression Mechanism

Fixed Evaluation Set

- A frozen set of 10–15 NSCLC patient vignettes
- No unknown fields
- Represents diverse biomarker and disease settings

Mandatory Re-Evaluation

Every update must be re-scored on the same vignette set using the same scoring sheets:

- Vignette Metadata Alignment
- Evidence Relevance
- Mapping Accuracy
- Citation Validity
- Misclassification flags

---

4. Non-Regression Criteria

Hard Stops (Update Blocked)

- Introduction of recommendation language
- New guideline or approval claims
- Increase in misclassification errors
- Removal of uncertainty disclosures

Soft Review (Allowed with Documentation)

- Changes in evidence relevance
- Shifts in mapping labels (e.g., aligned → silent)
- Addition of new evidence without loss of prior context

All soft changes must be documented. Each update produces a short update record describing what changed, why, and the observed non-regression impact.